

Every number is a rule applied to a sentence, and every rule leaves a trail back to the record

The engine's job is not to have opinions about 840 posts. It is to make each one auditable: tagged by fixed rules, joined to named groups, and reversible from any count on the site.

ONE RECORD, END TO END

1

RAW · REC-0302 · Reddit

"the insurance document. I remember roughly when it happened, but not the exact day. I searched a person's name and then browsed. After that I gave up. Sometimes I get a lot of plausible results and still cannot tell which one is right."

2

TAGGED BY RULE, ON SIX FIXED DIMENSIONS

object

document / screenshot

remembers

approximate time

lacks

time precision

state

exit path

outcome

abandoned

signals

5: strategy switch, browsing, abandonment ...

3

JOINED TO 10 NAMED GROUPS

need:N8 · need:N9 · seg:SEG-1 · node:D1:condition · claim:K01 · claim:K02 · claim:K05 · claim:K09 · hyp:H6 · hyp:H7

4

COUNTED, AND REVERSIBLE

This row is one of the 800 retrieval records, one of the 594 that name the identifier they lack, and one of the 341 that end on the exit path. Click any of those counts on the site and this record is in the list that opens.

Slides 3 and 4 are this same arithmetic run over all 800 rows: no step in between is a judgement call that a reader cannot check.

WHAT THE RULES REFUSE TO DO

An unstated outcome stays unknown, never inferred. Sentiment is never counted as behaviour. Source and platform are never used as segmentation axes. A sentence outside the 90 known types stops the run instead of being guessed. 37 tests fail the build if a headline count moves.

WHERE THE 840 ROWS COME FROM

Source	Rows	Relevant	Signals / row
Google Play Store	187	96%	1.78
Reddit	167	97%	1.77
Google Photos Community	152	95%	1.86
Apple App Store	117	97%	1.75
Social media	80	95%	1.84
YouTube comments	76	93%	1.82
Forums	61	87%	1.87

No platform tells a different story: relevance holds between 87% and 97%, and every source returns 1.75 to 1.87 effort signals per row. More scraping would buy nothing.

WHAT EACH AI LAYER MAY DO

- Rules (lexicon)** may classify and count coder of record; unknown text raises rather than guesses
- Model-assisted coder** may propose the same schema, quoting the record scored against the rules; its live run coded 0 of 100 rows, so it is not in use
- Grounded chat** may explain what the report already says may not add facts of its own or read the raw records